

# Introduction to Machine Learning

## Day 1

---

Machine Learning Course

Sammie Omranian

6/10/2024



# About instructor

---

- Sammie Omranian
- PhD student in Biomedical and Health Informatics
- Research area: Applications of Natural Language Processing in Healthcare
- Research Trainee at Brigham and Women's Hospital, Harvard Medical School

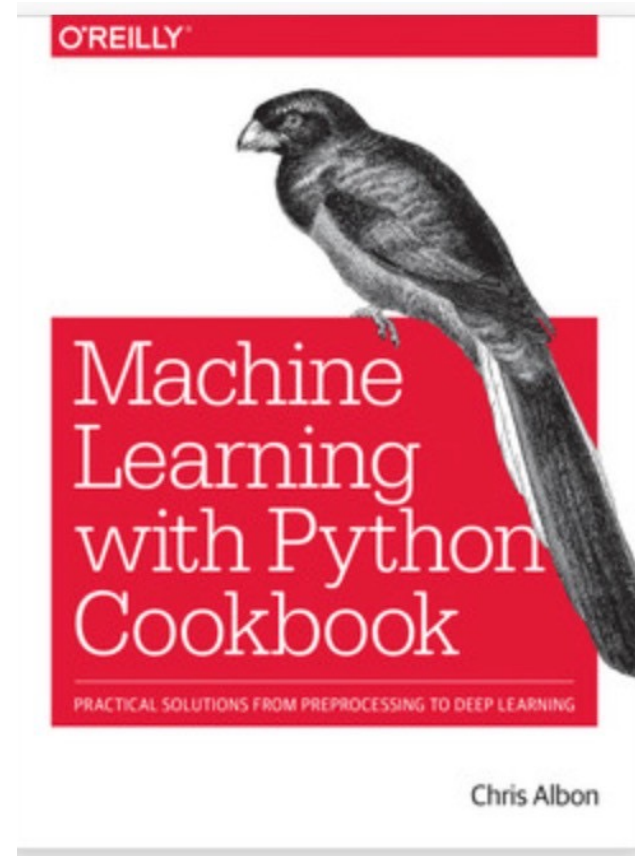


# Book

---

- Book description from [O'REILY](#) website:
- This practical guide provides nearly 200 self-contained recipes to help you solve machine learning challenges you may encounter in your daily work.
- Vectors, matrices, and arrays
- Handling numerical and categorical data, text, images, and dates and times
- Dimensionality reduction using feature extraction or feature selection
- Model evaluation and selection
- Linear and logical regression, trees and forests, and k-nearest neighbors
- Support vector machines (SVM), naïve Bayes, clustering, and neural networks

Saving and loading trained models



# Terminology Used in This Book

---

- Observation
  - A single unit in our level of observation—for example, a person, a sale, or a record.
- Learning algorithms
  - An algorithm used to learn the best parameters of a model—for example, linear regression, naive Bayes, or decision trees.
- Models
  - An output of a learning algorithm's training. Learning algorithms train models, which we then use to make predictions.
- Parameters
  - The weights or coefficients of a model learned through training.
- Hyperparameters
  - The settings of a learning algorithm that need to be set before training.



# Terminology Used in This Book

---

- Performance
  - A metric used to evaluate a model.
- Loss
  - A metric to maximize or minimize through training.
- Train
  - Applying a learning algorithm to data using numerical approaches like gradient descent.
- Fit
  - Applying a learning algorithm to data using analytical approaches.
- Data
  - A collection of observations.

# Introduction to Machine Learning

---

Arthur Samuel (1959), an American pioneer in the field of artificial intelligence, defined machine learning as:

*The field of study that gives computers the ability to learn without being explicitly programmed.*

He developed the checkers-playing program, in the late 1950s, one of his most notable achievements in the field of artificial intelligence and machine learning.



# Introduction to Machine Learning

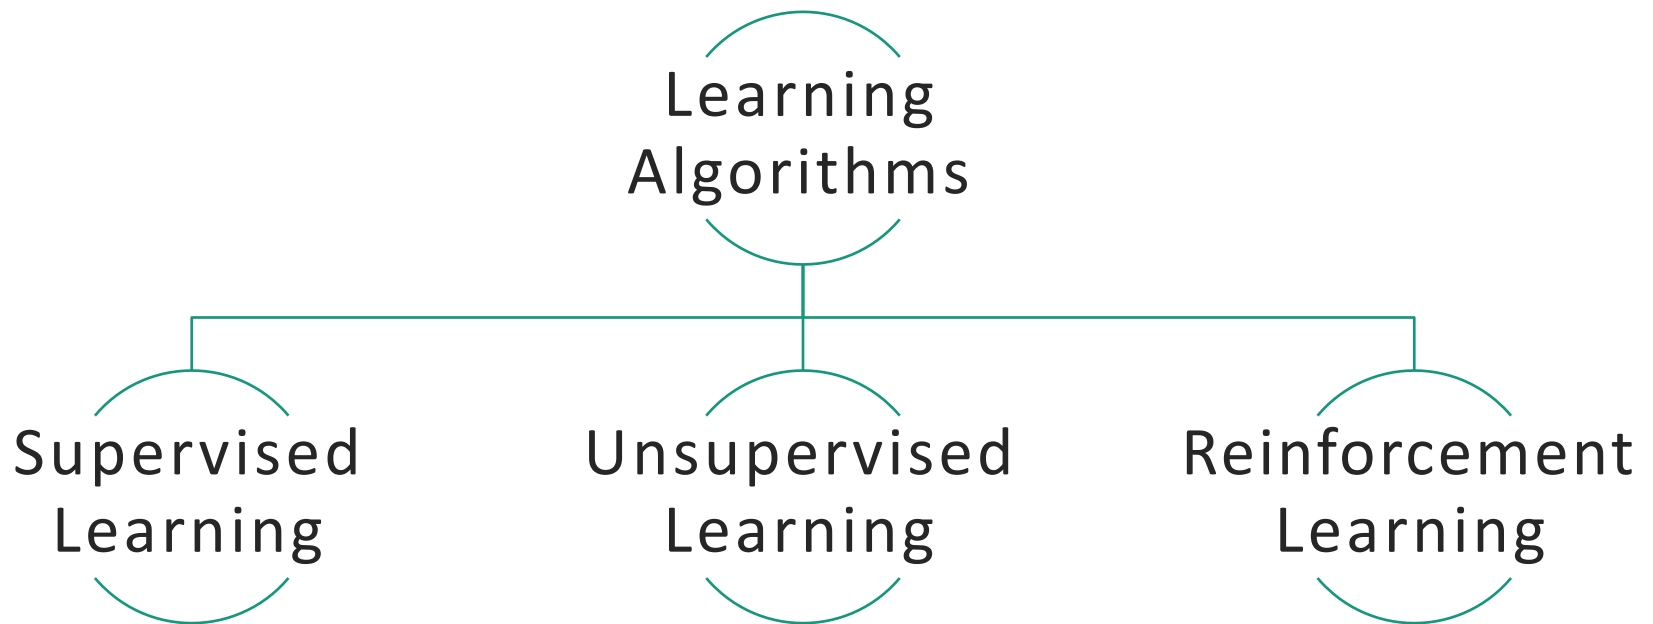
---

Tom Mitchell, a renowned computer scientist, defined machine learning as:

"A computer program is said to learn from experience  $E$  with respect to some task  $T$  and some performance measure  $P$  if its performance on  $T$ , as measured by  $P$ , improves with experience  $E$ ."



# Machine Learning Algorithms





# Introduction to Machine Learning

---

Types of ML:

- **Supervised Learning:** Learning from labeled data (e.g., classifying emails as spam or not spam).
- **Unsupervised Learning:** Finding patterns in unlabeled data (e.g., clustering customers based on purchasing behavior).
- **Reinforcement Learning:** Learning by interacting with an environment to maximize cumulative reward (e.g., training a robot to walk).



# Introduction to Machine Learning

---

- Applications:
  - Healthcare (predicting diseases)
  - Finance (fraud detection)
  - Marketing (customer segmentation)
  - Autonomous vehicles (self-driving cars)



# Introduction to Machine Learning

---

- Workflow:
  1. Data Collection
  2. Data Preprocessing
  3. Model Training
  4. Model Evaluation
  5. Model Deployment



# Supervised Learning

```
graph TD; A[Supervised Learning] --> B[Regression]; A --> C[Classification];
```

The diagram illustrates the structure of Supervised Learning. At the top, the text 'Supervised Learning' is centered. A horizontal line with a vertical drop in the middle connects it to two sub-categories below: 'Regression' on the left and 'Classification' on the right. Each category is enclosed within a green semi-circular arc.

Regression

Classification

# Supervised Learning

---

- **Regression vs Classification:**
  - Regression: Predicting continuous values (e.g., house prices).
  - Classification: Predicting categorical values (e.g., whether a tumor is benign or malignant).

# Supervised Learning

---

- Key Concepts:
  - **Features:** Independent variables used as input.
  - **Labels:** Dependent variables (target) we want to predict.
  - **Training Set:** Subset of data used to train the model.
  - **Test Set:** Subset of data used to evaluate the model's performance.

# Data Preprocessing

---

- Loading data with pandas

```
1 import pandas as pd  
2 df = pd.read_csv('data.csv')
```

- Handling missing values

```
1 df.fillna(df.mean(), inplace=True)  
2
```

- Feature scaling and normalization
- Splitting data into training and test sets

# Data Preprocessing

---

- Feature scaling and normalization



```
1 from sklearn.preprocessing import StandardScaler
2 scaler = StandardScaler()
3 df_scaled = scaler.fit_transform(df)
4 |
```

- Splitting data into training and test sets

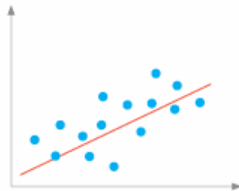


```
1 from sklearn.model_selection import train_test_split
2 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
3 |
```



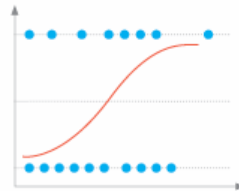
# Regression

## 5 types of regression



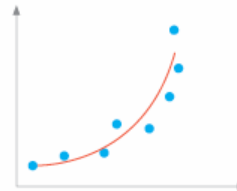
**Linear regression**

Predicts a continuous output by modeling a straight-line relationship between input features and target variables, such as estimating the impact of price changes on demand.



**Logistic regression**

Models the probability of binary outcomes, such as predicting customer churn; commonly used in classification tasks.



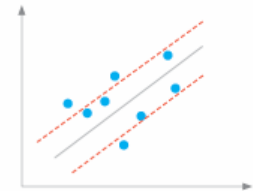
**Polynomial regression**

Captures nonlinear relationships, such as estimating the impact of ad spending on sales, by fitting a polynomial curve to data points.



**Time series regression**

Predicts future values in a time-dependent data set; often employed to forecast future values based on past observations, as seen in stock market analysis.



**Support vector regression**

Approximates a continuous function by identifying a hyperplane that best represents the data's structure; valuable in various applications, including financial market prediction.

# Linear Regression

---

- A linear approach to modeling the relationship between a dependent variable and independent variable.

- Mathematical Foundation:

$$y' = b + w_1x_1$$

- where:
    - $y'$  is the predicted value (a desired output).
    - $b$  is the bias (the y-intercept), sometimes referred to as  $w_0$ .
    - $w_1$  is the weight of feature 1. Weight is the same concept as the "slope" in the traditional equation of a line.
    - $x_1$  is a feature (a known input).

Although this model uses only one feature, a more sophisticated model might rely on multiple features, each having a separate weight:

$$y' = b + w_1x_1 + w_2x_2 + w_3x_3 + \dots + w_nx_n$$



# Linear Regression

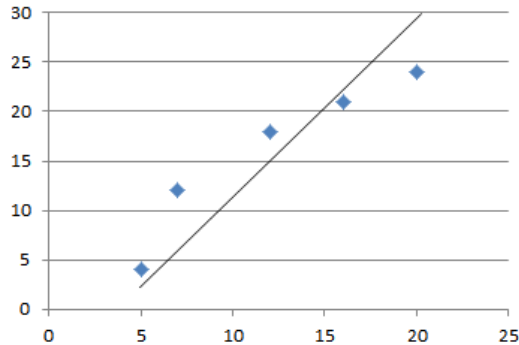


Diagram 1

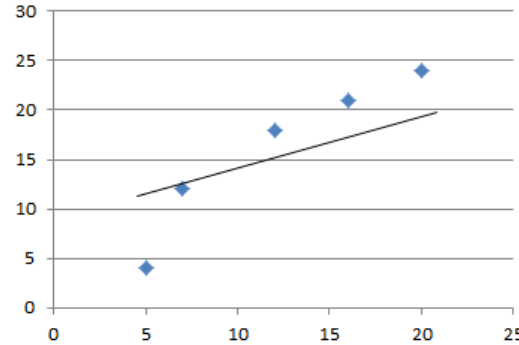


Diagram 2

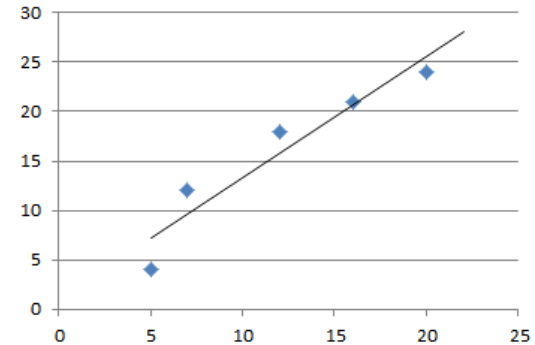
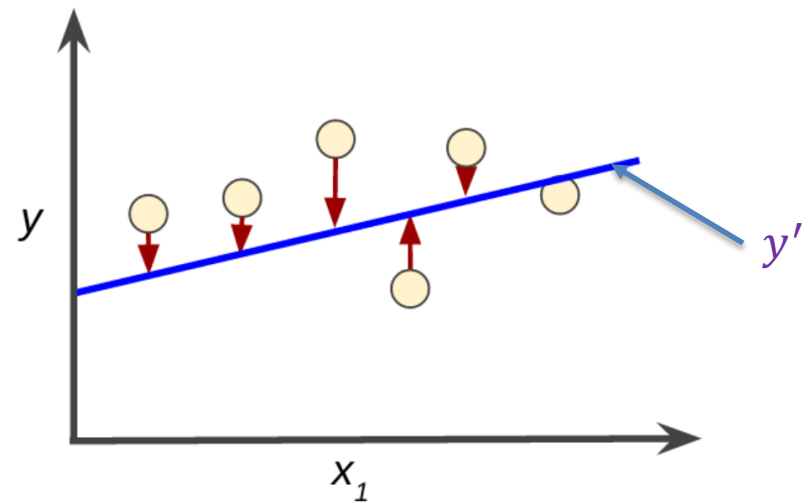
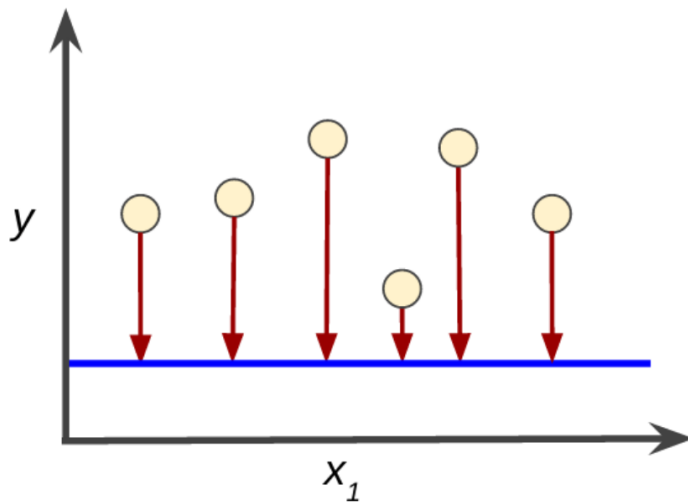
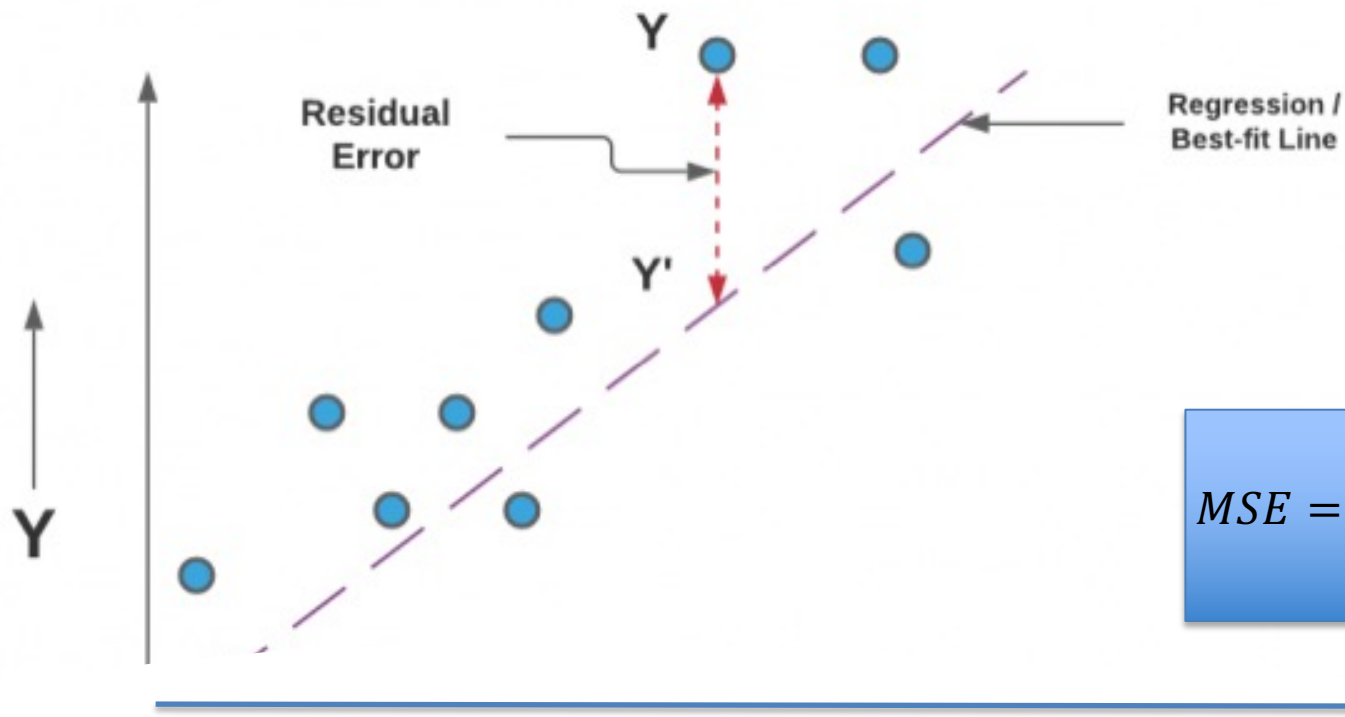


Diagram 3



High error (difference between actual value and predicted value) in the left model; low error in the right model.



$$MSE = \sum_{i=0}^n \frac{(y_i - y'_i)^2}{n}$$

X →

- **Cost Function:** Mean Squared Error (MSE)

The cost function defined as Mean Squared Error is the sum of the squared differences between the prediction and true value. And the output is a single number representing the cost. So, the line with the minimum cost function or MSE represents the relationship between X and Y in the best possible manner.

# Linear Regression

---

**Gradient Descent:** An iterative optimization algorithm to find values of the model parameters (bias and weights) that minimize the cost function.

Cost function  $\rightarrow$   $MSE = \sum_{i=0}^n \frac{(y_i - y'_i)^2}{n}$

Now , we can use gradient descent to find optimum parameters which is an iterative algorithm and apply the following rule “ update rule” .

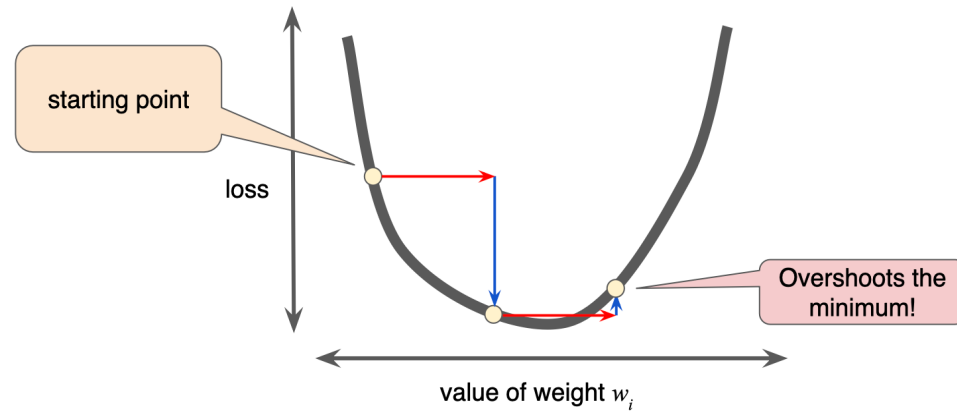
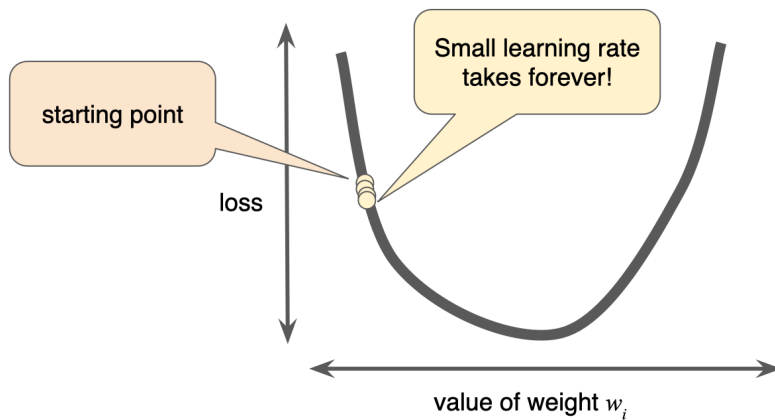
$$\begin{aligned} w_i &= w_{i-1} - \text{learning rate} * \partial(MSE(b, w_i))/\partial w \\ b_i &= b_{i-1} - \text{learning rate} * \partial(MSE(b_i, w))/\partial b \end{aligned}$$

gradient  $\rightarrow$

# Linear Regression

**Learning rate** decides the size of steps in gradient descent algorithm. The programmer sets this rate. It should not be too large nor too small. If it takes large steps, it may miss the optimum point. If it takes too small, it may take too many iterations and consume large computation time.

There's a Goldilocks learning rate for every regression problem. The Goldilocks value is related to how flat the loss function is. The common values are  $[0.00001, 0.0001, 0.001, 0.01, 0.1, 1]$ . You can check if your learning rate is doing well by plotting it on a graph.

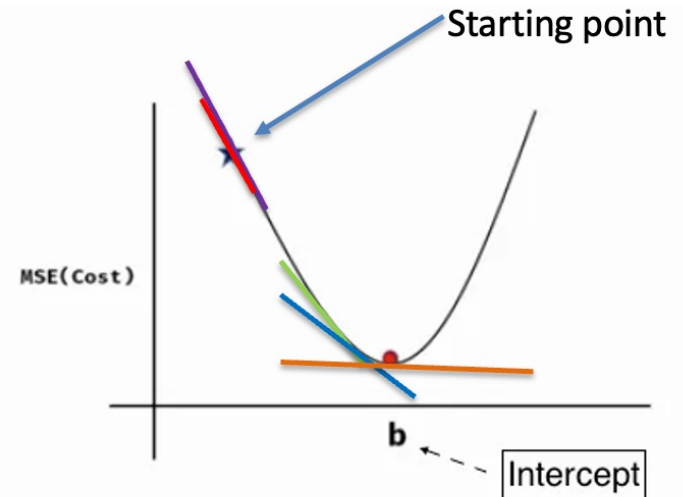
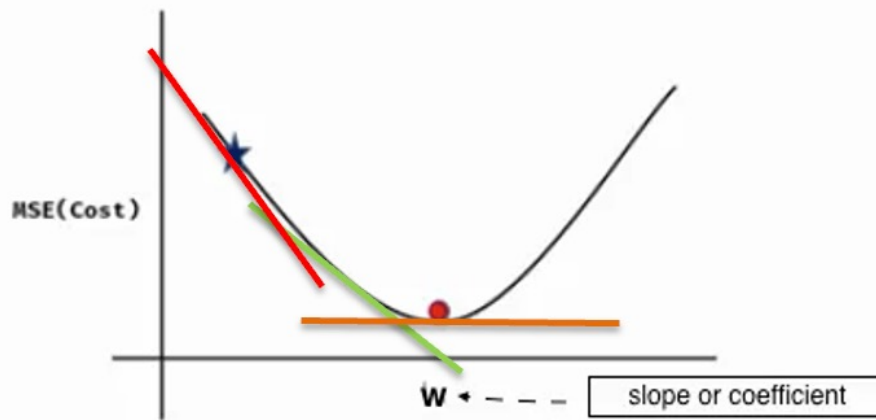


# Linear Regression

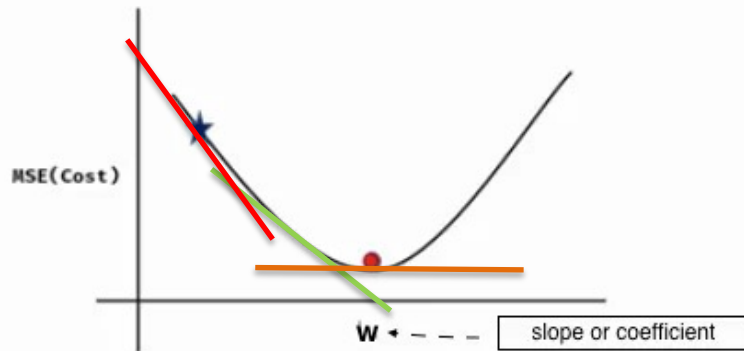
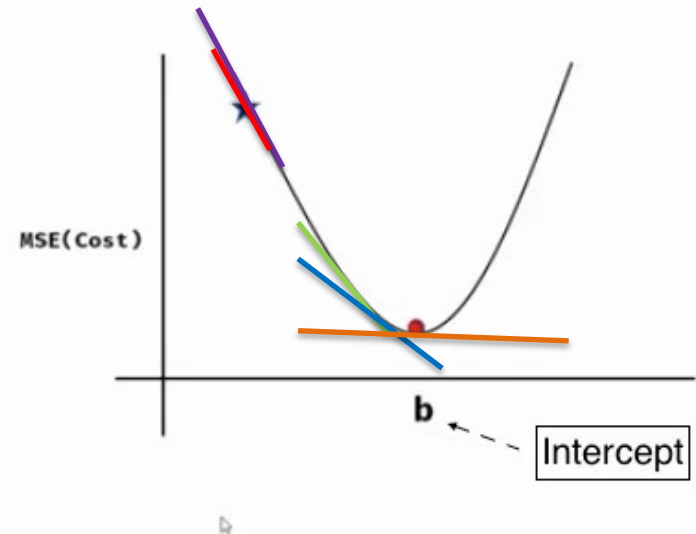
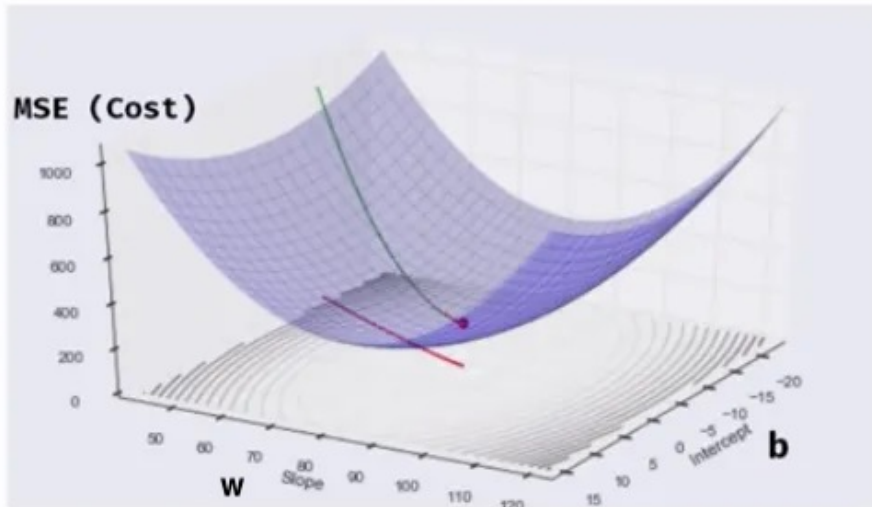
$$w_i = w_{i-1} - \text{learning rate} * \partial(MSE(b, w_i))/\partial w$$

$$b_i = b_{i-1} - \text{learning rate} * \partial(MSE(b_i, w))/\partial b$$

$$\text{precision} = \text{abs}(w_i - w_{i-1})$$



# Linear Regression



Gradient vectors:  $\frac{\partial(MSE)}{\partial w}$  ,  $\frac{\partial(MSE)}{\partial b}$

3D picture of gradient descent , MSE ( cost ) with weight (w) , and MSE with bias (b).

[Graph source](#)



# Linear Regression

---

- From Scratch:
- Step 1:
- Initialize Parameters:  $w_0, b_0, \text{learning rate}, \text{precision}, \text{\#iterations}$
- Step 2: Compute Gradient and update parameters
- Step 3: repeat until the difference between the values of parameters from two consecutive iterations is less than precision or when the number of iterations exceeds number of iterations, the algorithm should be stopped.

# Linear Regression

```
▶ 1 cur_x = 3 # The algorithm starts at x=3
  2 rate = 0.01 # Learning rate
  3 precision = 0.000001 #This tells us when to stop the algorithm
  4 previous_step_size = 1 #
```

```
▶ 1 while previous_step_size > precision and iterations < max_iters:
  2     prev_x = cur_x #Store current x value in prev_x
  3     cur_x = cur_x - rate * df(prev_x) #Grad descent
  4     previous_step_size = abs(cur_x - prev_x) #Change in x
  5     iterations = iterations+1 #iteration count
  6     print("Iteration",iterations,"\\nX value is",cur_x) #Print iterations
  7
  8 print("The local minimum occurs at", cur_x)
  9
```

# Regularized Regression

---

- In standard linear regression the model trains to minimize the sum of squared error between the true ( $y_i$ ) and prediction, ( $y'_i$ ) target values, or residual sum of squares (RSS):

$$RSS = \sum_{i=1}^n (y_i - y'_i)^2$$

- Regularized regression learners are similar, except they attempt to minimize RSS *and* some penalty for the total size of the coefficient values, called a shrinkage penalty because it attempts to “shrink” the model.

$$RSS + \text{penalty}$$

- There are two common types of regularized learners for linear regression:
  - Lasso regression
  - Ridge regressionThe only formal difference is the type of shrinkage penalty used.

# Lasso Regression

---

- **Lasso regression**—also known as L1 regularization—is a form of regularization for linear regression models.
- Lasso is similar to ridge, except the shrinkage penalty is a tuning hyperparameter multiplied by the sum of the absolute value of all coefficients:

$$\frac{1}{2n} RSS + \alpha \sum_{j=1}^p |\beta'_j|$$

- where  **$n$**  is the number of observations.
- **$\alpha$**  is the regularization hyperparameter, and  $\sum_{j=1}^p |\beta'_j|$  is the  $L_1$  penalty.

# Ridge Regression

---

- Ridge regression also known as L2 regularization
- In ridge regression, the shrinkage penalty is a tuning hyperparameter multiplied by the squared sum of all coefficients:

$$RSS + \alpha \sum_{j=1}^p (\beta'_j)^2$$

where  $\beta'_j$  is the coefficient of the  $j$  th of  $p$  features and  $\alpha$  is a hyperparameter.

- $\alpha$  is the regularization hyperparameter, and  $\sum_{j=1}^p (\beta'_j)^2$  is the  $L_2$  penalty.